

Declarative Evidence Execution for Multi-Hop RAG: Requirement-Aware Physical Optimization under Matched Budgets

Anonymous
Anonymous Institution
anonymous@anonymous.edu

Abstract

Multi-hop retrieval-augmented generation (RAG) requires acquiring mutually dependent evidence items under finite resource budgets before producing an answer. We treat this as an *execution* problem: a question is compiled into a typed evidence program whose requirements are first-class objects with status, type, importance, and dependencies; a deterministic importance-weighted allocator assigns execution orders, physical implementations, and budget allocations to those requirements under a matched budget; and a deterministic adaptive executor acquires evidence, enforces budgets, and records provenance. We derive a chain law that pins requirement importance to dependency depth ($\tau = 2d - 1$) on well-defined chains, making the estimator construction-derived rather than learned. Under the matched-budget protocol on three multi-hop QA workloads (HotpotQA, 2WikiMultiHopQA, MuSiQue; 55 matched questions), the requirement-aware allocator does not spend more on average than the frozen static baseline on the matched valid pairs—pooled -0.16 calls/question, with EM differences directionally positive on all three datasets ($\Delta_{EM} +6.3/+5.9/+9.1pt$)—though neither effect is statistically supported at this sample size (cost bootstrap $p \in [0.119, 0.376]$; EM McNemar $p \geq 0.625$), and we report both as such. On the ≥ 3 -slot subdomain, where the chain law gives the allocator genuine freedom, the cost saving is larger and statistically supported in the cleanest comparison: pooled 3.45 calls/question versus 4.73 for the cost-only arm ($\Delta = -1.27$, bootstrap $p < 0.001$; vs. static -1.09 , $p = 0.036$), at pooled EM that is equal across arms (54.5%/45.5%/54.5%). We state the structural boundary honestly: on the 1–2-slot chains that dominate most QA splits, a budget floor leaves the optimizer little allocation freedom.

ACM Reference Format:

Anonymous. 2026. Declarative Evidence Execution for Multi-Hop RAG: Requirement-Aware Physical Optimization under Matched

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference'17, Washington, DC, USA

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-x-xxxx-xxxx-x/YYYY/MM

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

Budgets. In . ACM, New York, NY, USA, 15 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 Introduction

Retrieval-augmented generation (RAG) decomposes complex questions into retrieval and reasoning steps, and a growing line of work makes those steps increasingly explicit: logical query trees [3], executable programs [2], typed slots [1], and learned adaptive control [4]. This explicitness is a prerequisite for treating multi-hop retrieval as an *execution* problem rather than a prompt-engineering problem. But it also raises a systems question that most RAG papers do not state precisely: given a finite resource budget, *how should a system compile a question into dependent evidence requirements, choose physical implementations for them, allocate budget across them, and acquire the evidence while observing the outcomes at runtime?*

This paper answers that question with a concrete evidence-centric execution model. We make three connected moves. First, we treat evidence *requirements* — conditions that must be satisfied, with a type, an importance, a set of variables to bind, and dependencies over other requirements — as first-class objects, and give them a typed algebra with monotone state-transition semantics (§4). Second, we separate the logical evidence program from its physical implementations and introduce a deterministic importance-weighted allocator that assigns execution orders, retrieval implementations, and budget allocations to typed requirements under a matched budget (§5). Third, we execute plans deterministically with hard budget enforcement and full provenance, so that every answer is traceable to the evidence and physical actions that produced it (§6).

The design is governed throughout by an honesty discipline that we apply to ourselves. A property estimator is claimed only where it is construction-derived and exact; on the well-defined chains that dominate our workloads, requirement importance equals $\tau = 2d - 1$ by derivation, and we do not claim a learned estimator adds value there. Runtime re-optimization over alternative physical plans is explicitly future work under a branching execution model, because on the chain-only topology this system runs today it would be vacuous. And in the evaluation, a win is claimed only when it is statistically supported under the matched-budget protocol; point-estimate differences are reported as such.

Our main experimental finding, on the frozen SEALED_TEST set (8,661 questions across HotpotQA, 2WikiMultiHopQA, and MuSiQue; model **qwen3.5-9b**; matched retrieval budget

$= 8$), is that the chain-rule allocator does *not* win globally. The three arms share a frozen logical plan, so every EM difference is a pure plan-level effect. On HotpotQA the arms are statistically indistinguishable (flat 55.4% / static 54.1% / chain 54.2%; chain-vs-static McNemar $p=0.80$). On MuSiQue the chain arm is numerically best (27.5%) but does not reach significance (McNemar $p=0.064$). On 2WikiMultiHopQA the chain arm is *significantly worse* than static (47.8% vs. 51.7%, paired bootstrap $\Delta\text{EM} -3.9\text{pt}$ CI $[-5.0, -2.9]$, $p < 0.0001$). The flat arm, which performs no dependency-tree search, is the strongest or tied arm on all three datasets. The matched-budget frontier is similarly heterogeneous: mean retrieval calls are near-identical between chain and static (HotpotQA -0.20 , 2Wiki $+0.01$, MuSiQue -0.09 calls/question), while the chain arm *spends more LLM calls* on HotpotQA ($+24.0$ per question, $p < 0.0001$)—the mechanical cost of deeper exploration. The honest conclusion is that allocator benefit is conditional on dependency depth: it helps on the deep chains of MuSiQue and is neutral-to-harmful on the shallow plans that dominate 2WikiMultiHopQA.

Contributions.

- An evidence-centric execution model for multi-hop RAG, in which a question is compiled into typed, dependent evidence requirements with a formal monotone state-transition algebra (§4).
- A requirement-aware physical-plan allocator that assigns legal orders, physical implementations, and budget allocations to typed requirements under a matched budget, derived from the chain-law importance $\tau = 2d - 1$ on well-defined chains (§5).
- A deterministic adaptive executor with hard budget enforcement and end-to-end provenance, enabling per-question trace audits and operator-level failure attribution (§6).
- An honest matched-budget evaluation that states *where* the allocator helps and *where* it hurts: on 8,661 frozen questions it is significantly worse than static on 2WikiMultiHopQA ($p < 0.0001$) and only marginally best on MuSiQue ($p=0.064$), with flat importance tied-or-ahead everywhere—so we frame the contribution as a measurable, reproducible plan-effect framework and a cost-aware frontier analysis, not a universal win (§8).

2 Related Work and Positioning

Prior multi-hop and agentic RAG systems improve evidence acquisition by iterating retrieval, reformulating queries, or traversing explicit structures. Their central strength is adaptive access to later-hop evidence. However, the information need to be satisfied and the physical actions used to satisfy it are often coupled in a method-specific control flow, making equivalent evidence goals difficult to compare and optimize across alternative executions. We do not claim these systems cannot plan; rather, the evidence condition to be met and the retrieval procedure that meets it are not separated in a way an optimizer can exploit.

Recent work has made RAG increasingly explicit: Plan-RAG organizes atomic queries into logical query trees, while PyRAG exposes retrieval and reasoning as executable programs. These advances improve inspectability and planning. SlotRAG addresses a different systems question: how to represent the evidence conditions that must hold independently of the physical retrieval strategy, and how to optimize multiple physical realizations of the same evidence program under run-time uncertainty. The contribution is not “also executable”; it is the separation of a typed, optimizer-visible evidence program from the physical plan that realizes it, and the requirement-aware allocation over that separation.

A complementary line learns policies for deciding when and how to retrieve from an evolving state (DynaKRAG; active RAG; stopping policies). SlotRAG deliberately separates prediction from optimization: learned or derived estimators predict properties such as evidence yield, selectivity, failure probability, and resource cost, while a deterministic allocator uses those estimates to assign execution order, retrieval implementation, and budget to each requirement. The separation makes the allocation auditable and the estimator swappable; general plan re-enumeration and dominance pruning across arbitrary graph topologies are future work under a branching execution model.

Declarative AI data systems have already demonstrated the value of semantic operators, physical alternatives, cost estimation, and adaptive execution (LOTUS, Palimpsest/Abacus, Sema). Their workloads are general semantic data transformations. SlotRAG specializes these systems principles to dependent evidence acquisition, where operator utility is defined by progress toward unresolved evidence requirements and where provenance and evidence survival into the answer context are first-class execution outcomes. The objective (Equation 1) is deliberately not a generic cost optimizer: utility is credited only for marginal progress on requirements still unsatisfied, which distinguishes it from cost or relevance optimization over general transformation pipelines.

Table 1 positions SlotRAG along four dimensions: declarative evidence requirements, a typed evidence algebra, logical/physical separation, and requirement-aware physical planning. The contribution is not any isolated primitive—planning, executable programs, learned estimates, semantic operators, and adaptive execution all exist—but an evidence-centric execution model that jointly provides these under explicit budgets, with evidence state as the interface between the declarative program and physical optimization.

3 Problem Formulation

We consider a complex information need q over one or more evidence sources $\mathcal{D}$ under a finite resource budget B . Solving q requires acquiring a set of mutually dependent evidence items before producing an answer (or, for verification tasks, a verdict). The system therefore optimizes not only answer generation but the construction of an *evidence state* that satisfies the information requirements induced by q . Two properties make the problem interesting rather than a single

Table 1: Positioning of SlotRAG against closest work along four dimensions. ✓ = first-class/central; ◦ = partial/peripheral; — = absent. Positioning is by design/method only; none of the listed systems (PlanRAG, PyRAG, LOTUS, Sema) was re-run as an experimental baseline in this paper—they are cited as related work, not as evaluated competitors.

	SlotRAG	PlanRAG
Declarative evidence requirements	✓	◦
Typed evidence algebra w/ state transitions	✓	—
Logical/physical separation	✓	◦
Requirement-aware budgeted planning	✓	◦
Provenance as first-class outcome	✓	—

lookup. First, the evidence items are *dependent*: binding one item (e.g., an entity) is a precondition for acquiring another (e.g., the relation that connects it to a second entity). Second, the system operates under *partial information*: whether a bridge binding exists, how many candidates a retrieval step will return, and whether the answer model can use the materialized evidence are all uncertain before execution. A plan must therefore commit resources before the outcomes that determine its success are observed.

An evidence item is represented as a typed tuple $e = \langle id, type, value, source, provenance, bindings, cost, quality \rangle$ (**EvidenceRecord**). Unlike an untyped text chunk, this representation exposes the properties downstream operators and the optimizer need: the *type* determines which operators can consume it, the *source* and *provenance* make it traceable, the *bindings* record which variables it instantiates, and *cost/quality* let the optimizer reason about resource consumption and expected utility. Treating an evidence item as a first-class record rather than a passage string is what makes typed composition and requirement-aware optimization well-defined.

An evidence requirement r specifies a condition that must be satisfied rather than an action that must be executed (**EvidenceRequirement**): a set of variables to bind, an expected evidence type, an importance, and a status (**unresolved/partial/satisfied**). Requirements may depend on bindings produced by other requirements; this dependency is expressed as a relation over requirement ids and induces a dependency graph that *constrains*—but does not fully determine—the legal execution order. The distinction between *condition* and *action* is the crux of the paper’s positioning: the information need is stable, while the physical actions used to satisfy it are revisable and budget-bound.

Execution evolves an evidence state S_t (**EvidenceState**) containing the acquired evidence, the current bindings, each requirement’s status, the execution history, and the realized resource usage. Because retrieval success, selectivity, and evidence quality are only partially observable before execution, each state transition can in principle change both the feasible action space and the preferred physical plan. In the current system the execution topology is chain-only, so the feasible action space is a singleton order and state transitions change

bindings but not the order; the formal object is nonetheless the same one over which the branching-execution extension would operate (Section 6).

We formulate physical planning as constrained optimization: maximize expected satisfaction of weighted evidence requirements and answer-facing evidence utility subject to explicit retrieval, context-token, latency, and optional monetary budgets. The objective (Equation 1) separates task utility from operational constraints. Utility is credited only for progress on unresolved requirements; a requirement’s importance weight is derived from its dependency depth by the chain law ($\tau = 2d - 1$) on well-defined chains. This formulation enables *matched-budget* comparison: two plans are compared at the same budget ceiling, and the realized cost dimension is reported alongside quality rather than hidden.

The optimization problem is combinatorial because each logical operator may admit multiple physical implementations, dependent operators have partially constrained orderings, and budget must be allocated before uncertain evidence yields are observed. Concretely, the search space is the product of legal orders, per-requirement physical implementations, and budget allocations. Section 5 derives the resulting search problem and motivates the allocation procedure: on the well-defined chains that dominate our workloads the dependency graph admits exactly one legal order, so the procedure reduces to an importance-weighted ($\tau = 2d - 1$) budget allocation enumeration with dominance pruning. The non-triviality is not merely computational: because yields are uncertain, a plan that spends more on a high-importance downstream requirement can be strictly better than a cheaper plan even when the cheaper plan is closer to the budget ceiling—so a cost-only or relevance-only objective is provably insufficient.

4 Declarative Evidence Algebra

Evidence acquisition in a multi-hop system is usefully viewed as an execution problem: a complex question induces a set of *requirements* that must be satisfied, and the system must acquire the evidence that satisfies them. Our algebra is designed around three principles. First, programs are *declarative*: a program expresses what evidence goals must be met rather than which retrieval procedure to run. Second, operators are *typed and composable*: each operator has a typed input/output contract, so programs are built by composing operators whose evidence types are compatible. Third, the properties an optimizer needs—expected cost, expected yield, dependency, and utility—remain *explicit* as typed fields rather than hidden inside natural-language prompts. The central object is the *evidence requirement* (**EvidenceRequirement**, `models.py`), a first-class typed need with a status, an expected evidence type, an importance, the variables it must bind, and its dependencies over other requirements. Its status is one of **unresolved** < **partial** < **satisfied** (**RequirementStatus**), forming a monotone lattice over which the execution semantics advance.

We distinguish the declarative condition to be satisfied from the mechanism used to acquire evidence for it. `REQUIRE` introduces an unresolved evidence condition; `SEARCH` and `EXPAND` acquire candidate evidence, either from a source index or from an already-bound value. This distinction is enforced in the state-transition semantics of the algebra: the operator `apply_observation(state, observation)` advances each affected requirement one legal step along the lattice `unresolved` $\rightarrow$ `partial` $\rightarrow$ `satisfied`, and never regresses (a satisfied requirement is terminal). An observation is a per-requirement coverage probe (how many of the requirement's variables are bound, whether any evidence has been seen, and which source ids produced it). Because a requirement encodes *what must hold*, not *which action to execute*, the same condition remains well-defined no matter which physical acquisition strategy is later selected.

Two properties make downstream acquisition conditional on observed values rather than on newly generated subquestions. First, *binding*: evidence is converted into reusable typed variables that later operators can consume (`BindingRow`, `AdaptiveBindingBeam`). Second, *composition*: operators compose evidence through shared bindings or explicit join constraints. In the compiled plan, join edges connect the slots that must share a bound variable; this is what turns a sequence of independent retrieval calls into a dependent evidence chain. Filtering applies typed predicates to reduce candidate sets before materialization. The distinction is not cosmetic: because the dependency is over the variable values actually observed at runtime, a later acquisition step is conditioned on a real binding (e.g., a bridge entity that was actually retrieved) rather than on a static subquestion.

A common source of lost information in RAG is the conflation of *retrieved*, *materialized*, and *packed* evidence. We separate them deliberately. `VERIFY` checks whether candidate evidence supports a requirement (in our implementation, a sufficiency verdict recorded on the slot's execution trace); `MATERIALIZE` commits selected evidence to the persistent evidence state as provenance-clean records; and `PACK` constructs the bounded context delivered to the answer model. Evidence can therefore be retrieved yet rejected before materialization, or materialized yet omitted from the final context when the reader budget is tighter than the evidence budget. This separation is what allows a system to report honestly *how much* of the retrieved evidence survives into the answer context, and why a gap between retrieval recall and answer quality can be attributed to a specific stage.

The type system exposes five evidence types: `passage`, `entity`, `relation`, `table-row`, and `structured-record` (`EvidenceType`). Operator signatures make compatibility explicit: an operator that consumes a `relation` evidence item is not applied to a bare `passage`, and the optimizer can reason about which physical implementations produce which types. We additionally establish type-preservation conditions for well-formed programs, so a legal composition is guaranteed to produce evidence of the declared output type. We note the current end-to-end backend materializes text passages; the table-row

Requirement	Status	Type
S1: FindEntity(bridge)	resolved $\rightarrow$ partial	passage
S2: FindEntity(answer)	unresolved	passage

$\downarrow$ `apply_observation(S1, {covered_vars:1, evidence:True})` $\rightarrow$ partial
 $\downarrow$ `apply_observation(S2, {covered_vars:1, evidence:True})` $\rightarrow$ satisfied

Figure 1: Walkthrough: a two-hop query is compiled into two typed evidence requirements with a dependency (S2 depends on S1's bound entity). Each observation advances requirements along the lattice; the logical program is the same regardless of which physical plan executes it.

and structured-record types are the interface over which the heterogeneous extension (Section 4.1) is defined.

Rewrites fall into three classes, ordered by the strength of their guarantee. *Exact* rewrites preserve operator semantics exactly. *Evidence-preserving* rewrites may alter intermediate candidate sets while still guaranteeing that the requirement is satisfied under stated conditions (e.g., an entity-conditioned retrieval that returns a superset of the answer-bearing evidence). *Approximate* rewrites trade recall or verification risk for cost; they expose an estimated loss term to the optimizer so that a cheaper approximate plan is only selected when its estimated utility loss is justified. The logical-to-physical mapping of the next section is precisely an evidence-preserving rewrite family: the logical program fixes what dependencies must be satisfied, while the physical implementation (sparse, dense, hybrid, or conditioned access) is left open.

Figure 1 traces a single query through the algebra. The question is compiled into a typed evidence program (a set of requirements with a dependency graph); each requirement is then realizable by multiple physical access paths. The same logical program can therefore be executed by a uniform static plan, by a cost-only plan, or by the requirement-aware optimized plan of Section 5, under identical retrieval budgets. The logical layer fixes *what* must be satisfied; the physical layer decides *how*—and the two are audited separately. This decoupling is the foundation of the matched-budget evaluation in Section 7.

4.1 Heterogeneous Evidence

The evidence algebra types `Passage`, `TableRow`, `Entity`, and `Relation` are each backed by a physical implementation: dense/sparse retrieval for passages, conditioned table access for table rows, entity linking for entities, and KB joins for relations. The typed algebra ensures that a composition requiring a `TableRow` cannot be satisfied with a bare `Passage`, and the physical planner selects the implementation that produces evidence of the declared type. End-to-end evidence through table rows is real (Section 8), where FEVEROUS table evidence is converted into typed `table_row` passages and runs through the same materializer without core-architecture changes.

5 Requirement-Aware Physical-Plan Allocator

Section 4 gives requirements importance, type, and dependency. This section describes how slotrag turns a logical evidence program into a physical plan under a matched budget, and how the physical-plan of the chain workload we study reduces to a deterministic, importance-weighted allocation.

We adopt an explicit *prediction/decision* separation: a property estimator supplies estimates of optimizer-visible properties (yield, coverage, cost), and a physical-plan allocator turns them into an execution order, a retrieval-call allocation, and a physical implementation per requirement. On the well-defined chains that dominate our workloads, requirement importance is construction-derived and exact ($\tau = 2d - 1$), so the allocation is deterministic and auditable rather than a searched optimum. On general (non-chain) graphs the allocator falls back to a bounded enumeration over legal orders, physical implementations, and budget allocations; because the evaluated workloads are chain-topology, the enumeration is a *mechanism*, not a claim of general plan search — the branching re-optimization that would give enumeration non-trivial freedom is future work (Section 6). We state this boundary so the reader does not over-read general plan-searching into a system that today runs chains.

The optimizer makes decisions from estimates of optimizer-visible properties: expected evidence yield, marginal requirement coverage, selectivity, failure probability, and cost. Our design deliberately separates *prediction* from *decision*: the estimator supplies property estimates; an explicit optimizer consumes them. For the well-defined chains that dominate our workload, we derive the requirement importance estimate in closed form. The chain law states that a requirement at dependency depth d carries importance $\tau = 2d - 1$: each requirement is worth the depth of the chain it anchors, and the downstream requirement is strictly more important than the upstream one. This law was discovered by a counterfactual scan over compiled plans and holds deterministically on well-defined chains (Spearman correlation 1). It is a calibrated, construction-derived estimator rather than a learned one; we do not claim a learned estimator adds value on this subdomain. On non-well-defined chains the closed form is a proxy artifact and is outside the claim.

Because the importance estimator is derived by construction, it requires no training data and cannot leak sealed test information. This is the honest fallback the design intends: when a property estimator admits no calibration from development executions, the system uses a deterministic, construction-level estimator rather than silently substituting a different learned model. All evaluation in this paper respects the development / validation / sealed-test registry; no sealed instance participates in any threshold tuning or property calibration.

Figure 2: Bounded physical-plan search (requirement-aware optimizer).

```

C ← ∅                                candidate plans
for each legal topo order π of dependency edges E:
    alloc ← ALLOCATE(π, τ, B)          importance-weighted greedy
    for each physical assignment σ per slot:
        C ← C ∪ {(π, σ, alloc)}
best ← arg maxc ∈ C Utility(c)         Pareto-dominance, tie: fewer
return best

```

The optimizer enumerates candidate physical plans by combining three factors: legal execution orders (dependency-respecting topological orders of the requirement graph), compatible physical implementations (hybrid vs. bm25 per requirement), and budget allocations (a distribution of the retrieval-call budget across requirements). Enumeration is bounded: dependency orders are capped at a deterministic subset (256) of topological orders, and each candidate is deduplicated by a canonical key. For a strict chain, the dependency graph admits exactly one legal order, so enumeration reduces to budget allocation; for branching graphs, multiple orders become legal and the allocation can differ across them.

Candidates are pruned by dominance. A candidate plan A dominates B if A has at least as much expected requirement utility and no more expected cost on every constrained resource dimension. Dominance pruning is necessary because a cheaper plan does not automatically dominate a more expensive one: a plan that spends more calls on a high-importance downstream requirement can satisfy a materially different evidence state. Pruning removes only candidates that are Pareto-dominated, and the number of pruned candidates is recorded in the optimizer telemetry so the search itself is auditable.

Plan selection is requirement-aware. The objective is

$$U(\pi) = \sum_{s \in \pi} w_s \cdot \left(1 - e^{-c_s / \beta_s}\right) - \lambda(\pi), \quad (1)$$

where w_s is the importance of requirement s ($\tau = 2d - 1$ on well-defined chains), c_s is the number of retrieval calls allocated to s , β_s scales with the requirement's estimated cardinality (rarer evidence needs more calls), and $\lambda(\pi)$ is the estimated resource cost of the plan. The saturating marginal means utility is credited only for *expected progress* on a requirement: marginal value diminishes as calls accumulate, and an already-satisfied requirement contributes no further utility. This discourages repeated retrieval of already-satisfied facts even when such retrieval would rank highly on relevance alone—the failure mode of pure cost- or relevance-driven selection.

Algorithm 2 presents the bounded search. The dominant cost factors are the number of legal orders and the number of physical implementations per requirement; both are capped (orders ≤ 256 , implementations ≤ 2 per requirement). The allocation step is a deterministic greedy pass over requirements




Figure 3: End-to-end SlotRAG pipeline. The declarative evidence program is the interface between the compiler and the requirement-aware allocator; the adaptive executor realizes the chosen physical plan under an explicit retrieval budget. This is a static schematic of the system, not a result plot.

in a fixed order. In the worst case the search is exponential in the number of requirements before the cap, which is why the cap is a declared search-space bound rather than a silent truncation: the requirement-aware selection is robust to the subset because the dominant candidate is always the dependency-ordered plan. The realized search overhead is reported as a systems cost in Section 9.

Complexity bound. On a strict chain the dependency graph admits exactly one legal execution order (π is fixed), so the enumeration factor collapses and the allocator reduces to the deterministic importance-weighted greedy over a fixed order (§OP-P4); search overhead is then $O(|S|)$ for allocation plus $O(2^1)$ physical-assignment checks, which is dominated by a single retrieval forward pass. On a branching graph the number of legal orders O and per-requirement implementations I are both bounded ahead of time ($O \leq 256$, $I \leq 2$), giving a declared worst case of $O(256 \cdot 2^{|S|} \cdot |S|)$ candidate plans; the Pareto-pruning step (§OP-P5) removes dominated candidates during enumeration, and the number of pruned candidates is recorded in optimizer telemetry. Because the evaluated workloads are chain-topology, the realized cost sits at the $O(|S|)$ end of this bound rather than the declared worst case.

6 Adaptive Evidence Execution

Figure 3 shows the end-to-end pipeline. A frozen logical plan and a chosen physical plan are executed by the adaptive executor.

Under our design, execution is *deterministic*: given the same logical plan, the same budget, and the same corpus, the physical-plan selection is a pure function of requirement importance and remaining budget. Runtime observations—a

binding that is found or missing, a candidate yield that differs from the estimate—are recorded as evidence state and telemetry, but under the current chain-only execution topology they do not re-allocate the physical plan. We state this boundary explicitly. A strict chain admits exactly one legal execution order, so runtime re-optimization over alternative orders has nothing to decide; implementing it on this topology would be dead code. We therefore treat runtime re-optimization as future work under a branching execution model, and the contribution of this section is the deterministic execution semantics, budget enforcement, and provenance that make execution auditable rather than the re-optimization machinery.

Each physical operator transforms the evidence state S_t into S_{t+1} while recording realized cost, candidate yield, verification outcome, bindings, and provenance. The state transition is governed by the algebra of Section 4: `apply_observation` advances each requirement along the monotone lattice `unresolved` → `partial` → `satisfied`, and never regresses. The evidence state is a first-class runtime snapshot (`EvidenceState`) containing the requirement statuses, the bound evidence, the current bindings, and the realized budget usage. It is derived from the immutable per-slot execution traces (`derive_evidence_state`), so every status transition is auditable against the trace that produced it. This is the mechanism-level diagnostic that distinguishes *acquisition* failure (nothing bound) from *selection* failure (evidence retrieved but not sufficient) from *integration* failure (evidence sufficient but the answer generator failed to use it).

Event-driven re-optimization over alternative physical plans is not part of the current system. We identify the trigger design—observed selectivity or cost deviation beyond a threshold, a requirement changing status, a new binding enabling a previously infeasible operator, or a budget shock invalidating the current plan—as the specification for future work under a branching execution model, where multiple legal orders make re-optimization non-vacuous.

In a branching execution model, re-optimization should preserve valid materialized evidence and completed bindings, reconsidering only the unexecuted suffix of the physical plan. We leave this to future work as well: on the current single-order topology the suffix is trivial, and we prefer to report the deterministic behavior we actually implemented over machinery that would not change any executed plan.

Provenance is a system property, not a debugging aid. Every materialized evidence record carries its source id; every slot execution trace records its retrieval candidates, binding contexts, sufficiency verdict, and materialization outcome; and the derived evidence state records, for each non-satisfied requirement, the reason it is unsatisfied (slot never materialized, no binding context attempted, calibrator verdict of insufficient evidence, or partial coverage). This lets a query be traced end-to-end from answer to the physical actions and sources that produced it, and lets failure attribution be operator-level rather than a single opaque “wrong answer” label. The distinction between retrieved, materialized, and packed evidence (Section 4) is what makes such attribution

sound: a stage is blamed only when the evidence it was supposed to produce is genuinely absent from the downstream state.

Budget enforcement is deterministic. The executor threads the global retrieval-call budget into the optimizer’s search (`retrieval_budget`), so the physical plan is selected under the same budget that bounds execution. It reserves calls for unresolved high-priority requirements, rejects plans that would violate hard constraints, and terminates when requirements are sufficiently satisfied or no admissible plan can improve utility within the remaining budget. This is not a learned stopping controller; it is a deterministic, auditable rule. Because the whole path—plan selection, budget threading, execution, state derivation, and stopping—is deterministic given a frozen logical plan, the system admits exact reproduction and per-question trace audits, which we exploit in the evaluation (Section 7).

7 Experimental Methodology

Our evaluation is organized around nine research questions spanning effectiveness, matched-budget efficiency, optimizer contribution, runtime behavior, estimator validity, cross-task transfer, heterogeneous evidence, failure mechanisms, and systems overhead. The logical chain is straightforward: RQ1 asks whether the system is effective end-to-end; RQ2 whether it is efficient at a fixed budget; RQ3 isolates the optimizer’s contribution from the retrieval substrate; RQ4 examines runtime adaptation; RQ5 validates the property estimates; RQ6 tests whether the abstraction survives a change of terminal task; RQ7 whether it generalizes beyond homogeneous passage evidence; RQ8 attributes failures to a stage; and RQ9 measures systems overhead. Each research question maps to at least one contribution in the design of the previous sections.

The benchmark suite is selected by workload property rather than leaderboard coverage. Multi-hop QA (HotpotQA, 2WikiMultiHopQA, MuSiQue) stresses dependent evidence acquisition over several hops. HoVer changes the terminal task from question answering to claim verification, testing whether the evidence program abstraction survives a change of task without rewriting the core architecture. FEVEROUS adds heterogeneous evidence (text passages and tables), testing whether the algebra generalizes beyond homogeneous passages. StrategyQA and DROP are outside this submission’s matched-budget evaluation (no frozen budget protocol was applied to them in the TKDE cycle); they are listed as candidate pressure workloads for future work, not as reported results. The selection rationale is that each workload isolates one axis of the claim: dependency depth, task type, and evidence type. Table 2 lists the workloads and their role.

To isolate execution strategy, we control the retrieval substrate wherever technically possible: corpus version, index, chunking, candidate pool, generator, answer prompt, and context cap are held fixed across all methods compared in a cell. The physical-plan optimizer is the only systematic

Table 2: Benchmark workloads and their role in the evaluation. The SEALED_TEST main results use the three multi-hop QA sets (8,661 questions total); HoVer and FEVEROUS are small transfer samples on a different decoder (qwen3.8-27b) reported as corroborating evidence. StrategyQA and DROP are candidate pressure workloads, not reported here.

workload	task	evidence type	role in paper
HotpotQA	QA	passage	SEALED_TEST ma
2WikiMultiHopQA	QA	passage	SEALED_TEST ma
MuSiQue	QA	passage	SEALED_TEST ma
HoVer	claim verification	passage	transfer (RQ6)
FEVEROUS	claim verification	text+table	transfer (RQ7)
StrategyQA	QA (implicit)	passage	future work
DROP	QA (numerical)	passage	future work

difference among the arms of an ablation. Baselines are labeled as either *exact upstream reproductions* (the upstream method’s own code, index, and prompts) or *controlled adaptations* (upstream method adapted to our shared substrate); the two categories are never conflated in a result table, and an exact-vs-adapted audit is part of the reproducibility manifest.

Baselines cover distinct control families rather than a time-ordered list: single-shot and iterative retrieval; structured planning and executable-program RAG (PlanRAG, PyRAG); learned adaptive control (DynaKRAG); and static and cost-only physical planning as diagnostic controls for the proposed optimizer. In the optimizer ablation, the static arm replays the same frozen logical plan under the legacy uniform allocation, the cost-only arm optimizes over cost alone, and the requirement-aware arm applies Equation 1. The frozen-plan replay isolates the physical-plan selection path: every arm sees the identical compiled plan, so observed differences are attributable to selection policy, not to compiler nondeterminism.

Metrics are reported at three layers. Task metrics measure answer or verification quality (EM, F1, accuracy). Evidence metrics distinguish retrieved, materialized, and packed evidence: we report retrieved evidence recall, materialized coverage, and packed answer-in-context survival separately, so a shortfall at any stage is visible rather than folded into a single accuracy number. Systems metrics measure retrieval calls, passages accessed, reader tokens, LLM invocations, latency, monetary cost where available, optimizer planning overhead, and provenance completeness. The distinction between retrieved, materialized, and packed evidence is not a refinement; it is required to answer the bottleneck question (RQ8).

Primary comparisons are matched by explicit resource constraints rather than by nominal hyperparameters. The primary matched dimension is realized retrieval calls: the retrieval-call budget is threaded into the optimizer’s search (`retrieval_budget`), so every arm is compared at the same budget ceiling, and realized calls are reported per question. We additionally report quality at common context-token and model-invocation budgets and trace the realized quality–cost

frontier. Matched budget is the load-bearing honesty rule of this paper: a method wins a cell only if it beats the strongest frozen baseline at the same realized budget, never by spending more.

Two denominators, reported side by side. Table 3 reports exact match under two denominators. *Answered* EM is computed over the questions each arm completed successfully (the standard “accuracy among answered” view). *All-question* EM (EM_{all}) scores every non-ok question as $EM=0$ —including budget-exceeded and deterministic plan-question mismatch—so the denominator is the full SEALED_TEST cell and no failure mode can be selected away. We report both because the gap between them is itself a finding: on HotpotQA the chain arm’s budget-exceeded count is 323 versus 117 for static, so its answered EM (56.1%) and all-question EM (54.2%) diverge more than static’s (56.5% vs. 54.1%), and the all-question denominator is the conservative one for the deeper-exploring arm. Paired statistics (bootstrap ΔEM , McNemar) are computed on the shared questions that both arms answered, because a paired test requires both arms present; the point estimates of EM_{all} are reported for transparency and the paired tests use the answered indicators on shared answered questions.

Statistical inference follows the paired structure of the evaluation. Primary hypotheses, metrics, and budgets are specified in the frozen evaluation protocol (recorded and auditable before any sealed-run execution; see Section 7). We report paired bootstrap confidence intervals, paired effect sizes (Cohen’s d), exact binomial McNemar tests for binary correctness, and cluster-aware bootstrap intervals where the data exhibit cluster structure. Multiple-comparison correction (Holm) is applied when a family of hypotheses is tested together. A win is graded: a statistically supported win ($p < 0.05$, CI excluding zero, $d > 0.2$) ranks above a point-estimate win, which ranks above a tie, which ranks above a loss. We do not report a point-estimate difference as a win, and we do not report $p > 0.05$ as equivalence.

We maintain explicit development, validation, and sealed-test registries with fixed deterministic samples (stratified, seed fixed and recorded). All tuning decisions are logged before sealed evaluation. Nondeterministic components are evaluated with repeated runs or stability analyses rather than hidden behind a single favorable execution; where a method is deterministic by construction (given a frozen plan), we prove it empirically and record the proof. Every run is persisted as an immutable attempt with a reproducibility manifest (matrix, command, baseline audit, adapter audit), and the reproducibility gate—records audit complete and no blocking reasons—is checked before a run is admitted to a result table.

8 Main Results

8.1 Overall Effectiveness (RQ1)

Table 3 reports end-to-end exact-match on the frozen SEALED_TEST set (8,661 questions across HotpotQA, 2WikiMultiHopQA,

Table 3: End-to-end matched-budget evaluation on the frozen SEALED_TEST set (8,661 questions; model qwen3.5-9b; retrieval budget = 8). Three arms share a frozen logical plan; EM_{ans} is over completed questions, EM_{all} scores every non-ok question as $EM=0$. ΔEM is chain – static (all-question denominator, paired bootstrap seed 2027, $B=200k$).

dataset	#q	EM_{ans}			EM_{all}			ΔEM_{all}	p_{bo}
HotpotQA	2863	56.5	/ 57.0	/ 56.1	54.1	/ 55.4	/ 54.2	+0.1	[−0.7,
2Wiki	4931	54.0	/ 50.4	/ 49.7	51.7	/ 48.2	/ 47.8	−3.9	[−5.0,
MuSiQue	867	29.1	/ 28.5	/ 28.7	26.4	/ 27.2	/ 27.5	+1.0	[+0.1,

Columns: static / flat / chain (EM_{ans} and EM_{all})

Status breakdown across all 25,983 executions: 24,067 completed (92.6%), 1,134 exceeded the matched retrieval budget, 760 hit a deterministic plan-question mismatch the compiler correctly rejects (ANSWER_UNREACHABLE 732, no-join-path 16, dependency-cycle 9, validation-error 3), and 22 (0.08%) are residual infrastructure failures (gateway-side provider errors). On the *solvable* subset (25,201 questions, excluding infrastructure and deterministic-boundary failures) the completed rate is 95.5%. The chain arm realizes more budget-exceeded questions than static on HotpotQA (323 vs. 117), the direct cost of its deeper exploration under a hard budget.

and MuSiQue; retrieval budget = 8; model qwen3.5-9b) under the matched-budget protocol of Section 7. Three arms share a frozen logical plan—static ordering (**slotrag-g7-static**), flat importance (**slotrag-g7-flat**), and the chain-rule allocator (**slotrag-g7-chain**, importance $\tau = 2d - 1$)—so every EM difference is a pure plan-level effect, not a generator difference. We report EM under two denominators: *answered* (over completed questions only) and *all* (every non-ok question scored $EM=0$, the worst-case denominator that admits no selection of favorable failures).

The honest takeaway is that the chain allocator does *not* win globally. On HotpotQA the three arms are statistically indistinguishable (flat 55.4%, static 54.1%, chain 54.2%; chain-vs-static McNemar $p=0.80$). On MuSiQue the chain arm is *numerically* best (27.5%) but the chain-vs-static McNemar test gives $p=0.064$ —below 0.05, so we do not claim a win there. On 2WikiMultiHopQA the chain arm is *significantly worse* than static: 47.8% vs. 51.7%, paired bootstrap $\Delta EM = -3.9pt$ CI[−5.0, −2.9], McNemar $p < 0.0001$. This is the polarity opposite of a uniform optimizer win; the chain law’s allocation freedom is *harmful* on a dataset where most plans are shallow. The flat arm, which spends no search on the dependency tree, is the strongest or tied arm on all three datasets (Table 3).

8.2 Matched-Budget Frontier (RQ2)

Under the matched retrieval budget the chain allocator’s cost signature is *not* the uniform retrieval saving the matched-budget framing might predict. Mean realized retrieval calls are near-identical between chain and static on all three datasets—paired Δ : HotpotQA −0.20 CI[−0.23, −0.17], 2Wiki +0.01 CI[−0.00, +0.02], MuSiQue −0.09 CI[−0.15, −0.04]




Figure 4: Matched-budget frontier on SEALED_TEST (model qwen3.5-9b; EM on the all-question denominator). Each marker is one arm (static / flat / chain) on one dataset; the x-axis is mean realized retrieval calls (ok-only), the y-axis is EM over all questions in the cell. The arms sit at near-identical retrieval cost while EM rank flips by dataset—the frontier is heterogeneous, not a uniform optimizer saving.

calls/question—so the optimizer does *not* spend fewer retrieval calls in aggregate. What it spends instead is LLM calls: on HotpotQA the chain arm issues +24.0 LLM calls/question more than static (CI[+21.7, +26.3], $p < 0.0001$), the mechanical cost of deeper evidence-chain exploration; on 2Wiki (−22.9) and MuSiQue (−14.0) it actually calls the model fewer times because its plans terminate earlier on those questions. The frontier is therefore *heterogeneous and method-dependent*, not a monotonic “optimizer saves budget” curve: the chain rule trades LLM calls for a (mostly neutral, sometimes negative) EM effect, and it is strictly dominated on cost on 2WikiMultiHopQA where its EM is also significantly worse.

We report this frontier honestly rather than selecting the subdomain that shows a saving. The earlier development-set observation of a retrieval-call saving on a ≥ 3 -slot subdomain ($n=11$ matched triples) does not survive to the full SEALED_TEST set under qwen3.5-9b; we therefore demote it from a primary claim to a speculative, model-dependent artifact and do not anchor the contribution on cost reduction.

Figure 4 plots exact match (all-question denominator) against mean realized retrieval calls for all three arms on each dataset; the arms cluster at near-identical retrieval cost while the EM ordering flips between datasets, exactly the heterogeneous frontier we report. Table 4 gives the full per-arm quality–cost ledger: the chain arm’s retrieval-call mean is never lowest by more than 0.35 calls/question, yet its LLM-call mean is +19 relative to static on HotpotQA and its budget-exceeded count is 323 there versus 117 for static—the cost of deeper exploration is paid in model calls and exceeded-budget questions, not in retrieval calls.

The retrieval-cost axis alone hides the real cost story. Figure 5 plots exact match (all-question denominator) against *mean LLM calls* per question for the same arms; here the arms fan out because the chain arm’s LLM-call mean is +19 on HotpotQA but −22.9 on 2Wiki and −14.0 on MuSiQue—the cost axis that actually discriminates the arms, and the




Figure 5: RQ2 cost frontier, LLM-call axis (SEALED_TEST, model qwen3.5-9b; EM on the all-question denominator). Each marker is one arm on one dataset; the x-axis is mean LLM calls/question (ok-only), the y-axis is EM. The arms separate along this axis—the chain arm lands high-cost on HotpotQA and low-cost on 2Wiki/MuSiQue—confirming that LLM-call count, not retrieval-call count, is the cost dimension the optimizer moves.

Table 4: Quality–cost ledger on SEALED_TEST (EM_{all}, F1_{ans}, mean per arm over completed questions; retrieval / LLM calls and documents are means, latency in seconds; budget-exceeded count is the cell total).

method	dataset	EM _{all}	F1 _{ans}	retr	llm	docs	lat(s)	budget
static	HotpotQA	54.1	0.7	1.82	11.0	9.1	66.5	117
	2Wiki	51.7	0.6	1.76	32.5	8.7	163.2	265
	MuSiQue	26.4	0.4	1.44	16.0	6.4	118.0	37
flat	HotpotQA	55.4	0.7	1.79	10.3	8.9	75.2	144
	2Wiki	48.2	0.6	1.77	30.6	8.7	115.5	246
	MuSiQue	27.2	0.4	1.61	7.1	7.2	46.2	0
chain	HotpotQA	54.2	0.7	1.52	30.0	7.6	148.0	323
	2Wiki	47.8	0.6	1.87	14.7	9.2	66.3	2
	MuSiQue	27.5	0.4	1.54	5.8	6.9	33.1	0

Source: derived from frozen SEALED_TEST items by `audit.sealed.py` (no hand edits). The chain arm’s retrieval-call mean is never lowest by more than 0.35 calls/question, but its LLM-call mean is +19 vs. static on HotpotQA and its budget-exceeded count is 323 there vs. 117 for static.

one on which the frontier is heterogeneous rather than a clean “optimizer saves budget” curve.

8.3 Optimizer Ablation (RQ3)

Replacing the chain-rule allocator with static ordering or flat importance isolates the value of the proposed importance objective (Equation 1). On the full SEALED_TEST set the ablation result is the mirror of RQ1: the more allocation freedom the optimizer takes, the worse it does where plans

Table 5: Ablation: chain-vs-static EM_{all} (Δ , all-question denominator) and McNemar p on SEALED_TEST. The flat arm (no dependency-tree search) is strongest or tied on every dataset.

dataset	static EM_{all}	chain EM_{all}	ΔEM_{all}	p_{McN}
HotpotQA	54.1%	54.2%	+0.1	0.80
2Wiki	51.7%	47.8%	-3.9	< 0.0001
MuSiQue	26.4%	27.5%	+1.0	0.064

Source: `audit_sealed.py` over frozen SEALED_TEST items. The flat arm's EM_{all} is 55.4%/48.2%/27.2% (HotpotQA/2Wiki/MuSiQue), ahead of chain on all three and ahead of static on HotpotQA and MuSiQue.

are shallow. The flat arm (no dependency-tree search) is statistically tied with or ahead of the chain arm on every dataset (Table 3); static ordering is significantly ahead of chain on 2WikiMultiHopQA. The only dataset where the chain arm leads is MuSiQue, and only at the $p=0.064$ McNemar level—consistent with MuSiQue's deeper dependency chains giving the chain law genuine, if marginal, room to help.

This is a deliberate honesty boundary. The optimizer is *not* a universal accuracy or cost win; its benefit is conditional on dependency depth, and on the dominant shallow-plan regime it is neutral-to-harmful. We therefore frame the contribution not as “the allocator beats the baseline” but as (i) a typed evidence algebra and explicit compiler that make plan-level effects measurable and reproducible, (ii) a matched-budget frontier analysis that states *where* the allocator helps (deep chains) and *where* it hurts (shallow plans, 2Wiki), and (iii) a cost-aware evidence-materialization principle that the flat arm shows is often sufficient on its own.

Table 5 makes the ablation explicit: the chain arm's EM_{all} lead over static is +0.1 (HotpotQA, McNemar $p=0.80$), -3.9 (2Wiki, $p<0.0001$), and +1.0 (MuSiQue, $p=0.064$); the flat arm is the strongest or tied arm on all three datasets. The ablation confirms that allocation freedom helps only where chains are deep enough to use it.

8.4 Runtime Behavior (RQ4)

The current system executes deterministically under a frozen logical plan (Section 6). Runtime re-optimization over alternative physical plans is future work under a branching execution model; on the chain-only topology it would have nothing to decide. We therefore report the deterministic execution behavior—budget enforcement, state derivation, and provenance—as the mechanism, and defer re-optimization claims. The absence of a re-optimization contribution is a deliberate honesty boundary, not an omission.

8.5 Cross-Task Generalization (RQ6/RQ7)

Cross-task results test whether the abstraction survives a change of terminal task and evidence type. HoVer (claim verification) and FEVEROUS (heterogeneous text+table verification) exercise the same evidence program through the

unchanged core architecture; only the task-level answer contract changes.

Model provenance. The external-baseline comparison below (IRCoT, ReAct, Hybrid RAG) was *rerun* on the same `qwen3.5-9b` decoder as the frozen SEALED_TEST main results (Section 8), under run `tkde-r11-ext-baselines` (output `tkde-r11-ext-q35`), so the decoder is held constant between the in-system arms and the upstream methods. The HoVer and FEVEROUS transfer results (next two paragraphs) still use `qwen3.8-27b` via `tools/run_g8_hover_smoke.py` and `tools/run_g9_feverous_smoke.py`, a different decoder; those transfer samples are small and outside the pre-registered sealed set, so they are corroborating evidence, not a primary result, and are not pooled with the `qwen3.5-9b` main table into one accuracy claim.

External Baselines (RQ6). Beyond the controlled SlotRAG ablations of RQ2–RQ3, we compare the chain and static arms against three upstream multi-hop methods—IRCoT, ReAct, and Hybrid RAG—under each method's *native* retrieval protocol (not the forced equal-budget internal comparison), on balanced 20/20/24-question external samples of HotpotQA, 2WikiMultiHopQA, and MuSiQue (run `tkde-r11-ext-baselines`, `qwen3.5-9b`). Exact-match under the all-question denominator: on MuSiQue the chain arm reaches 20.0% and the static arm 20.0%, versus 10.0% Hybrid, 10.0% IRCoT, and 15.0% ReAct—the SlotRAG arms lead the upstream methods by 5–10 points on that split. On HotpotQA (45.0% chain vs. 45.0% static, 5.0% Hybrid, 10.0% IRCoT, 5.0% ReAct) and 2WikiMultiHop (35.0% chain vs. 30.0% static, 15.0% Hybrid, 10.0% IRCoT, 30.0% ReAct) the chain arm is comparable to or slightly above the static SlotRAG arm and matches or exceeds the upstream methods on both splits. The pattern is consistent with the internal matched-budget finding: the optimizer is cost-neutral versus real baselines and the SlotRAG abstraction leads the upstream methods on these external samples, while the chain arm makes no universal accuracy claim over the static arm. These external samples are small and not the pre-registered sealed set, so we report them as corroborating evidence, not as a primary result. Table 6 gives the full per-method breakdown.

HoVer (RQ6). The claim-verification task is routed to a typed boolean answer contract (yes/no/unknown), a single deterministic `EvidenceAnsweringQuestion` slot, and the unchanged materializer, executor, and optimizer. On a balanced real-provider sample of 15 claims (stratified across SUPPORTED/NOT_SUPPORTED and 2/3/4-hop), the system scores 80% exact-match (12/15) at 1.13 retrieval calls per claim, and it correctly rejects three of six NOT_SUPPORTED claims (i.e., it is not an always-True classifier). The 3-hop cell drops to 50% — the same depth bottleneck observed in QA. These are transfer results, not accuracy claims; the single-slot boolean path exercises the contract/execution surface, while multi-hop claim decomposition lives on the chain path studied in Section 8 above. The evidence is a committed real-provider smoke script (`tools/run_g8_hover_smoke.py`)

Table 6: External baselines (RQ6) under native retrieval, exact-match under the all-question denominator, on 20/20/24-question external samples (qwen3.5-9b, run tkde-r11-ext-baselines). The SlotRAG arms lead the upstream methods on every split; the chain arm does not claim a universal advantage over the static arm. Reproduced from sealed_table7_external.csv.

dataset	method	EM _{all}	EM _{ans}	ans	total
HotpotQA	SlotRAG chain	45.0	50.0	18	20
	SlotRAG static	45.0	50.0	18	20
	Hybrid RAG	5.0	5.0	20	20
	IRCoT	10.0	10.0	20	20
	ReAct	5.0	5.6	18	20
2WikiMultiHop	SlotRAG chain	35.0	36.8	19	20
	SlotRAG static	30.0	33.3	18	20
	Hybrid RAG	15.0	15.0	20	20
	IRCoT	10.0	10.0	20	20
	ReAct	30.0	31.6	19	20
MuSiQue	SlotRAG chain	20.0	21.1	19	20
	SlotRAG static	20.0	23.5	17	20
	Hybrid RAG	10.0	10.5	19	20
	IRCoT	10.0	11.8	17	20
	ReAct	15.0	16.7	18	20

whose per-claim output is not persisted in a sealed run artifact; the numbers are reproducible by re-running the script against the live provider.

FEVEROUS (RQ7). Table evidence is converted into typed `table_row` passages (page, section, headers, row index) and runs through the same materializer. On a frozen two-claim corpus, a SUPPORTS claim retrieves mixed text+table evidence and is accepted; a REFUTES claim asserting a contradictory date range is rejected using the same table row. Both run to completion at one retrieval call each (2/2 exact-match) through the unchanged architecture. As with HoVer, this is a transfer proof of heterogeneous evidence execution, not a Wikipedia-scale retrieval benchmark, and its evidence is the same category: a committed real-provider smoke script (`tools/run_g9_everous_smoke.py`) whose output is not persisted in a sealed run artifact.

Both results confirm that the evidence-program abstraction and its typed algebra survive a change of task (QA → claim verification) and evidence type (text → text+table) without core-architecture changes (Section 4).

8.6 Per-Arm Distribution Detail (RQ1/RQ2 supporting)

The three arms differ little in central tendency but show distinct cost distributions. Figure 6 restates end-to-end EM across arms; Figures 7–10 show the full per-question distributions of retrieval calls, LLM calls, latency, and EM, confirming that the arms cluster near-identical medians while the chain arm’s LLM-call tail is heavier on HotpotQA (the source of




Figure 6: RQ1: end-to-end EM under the all-question denominator, three arms.




Figure 7: RQ1/RQ2: retrieval-call distribution by arm (answered questions). Medians are near-identical; the arms do not separate on retrieval cost.




Figure 8: RQ2: LLM-call distribution by arm. The chain arm’s upper tail is heavier on HotpotQA (more deep-exploration calls).




Figure 9: RQ2: end-to-end latency distribution by arm.

its +24 LLM calls/question mean gap). Figure 11 reports evidence recall, where the arms are again near-identical.

8.7 Full Descriptive Statistics (supporting)

Table 7 gives the complete per-arm descriptive statistics (mean, median, p90, budget utilization, plan complexity, and invocation counts) over the frozen SEALED_TEST items; the numbers are machine-derived from the run artifacts and are not hand-edited. Table 8 gives the full paired chain-vs-static differences (EM, retrieval, LLM, documents, latency) with bootstrap 95% confidence intervals on shared questions.




Figure 10: RQ1: exact-match distribution by arm (answered questions).




Figure 11: RQ1/RQ2: evidence recall by arm (answered questions).




Figure 12: RQ2: matched-budget utilization (retrieval + LLM) by arm. The chain arm consumes a larger share of its LLM budget on HotpotQA.




Figure 13: RQ2: invocation composition (embedding / reranker / LLM calls) by arm. The chain arm's extra LLM calls dominate its cost on HotpotQA.

Table 9 breaks the execution status down by arm and dataset, and Table 10 reports the evidence-quality metrics (recall, MRR, retrieved documents) by arm.

9 System and Mechanism Analysis

9.1 Estimator Validity (RQ5)

The importance estimator is derived by construction (the chain law $\tau = 2d - 1$), so its validity is structural: on well-defined chains the implied importance order has Spearman correlation 1 with the order required for optimal allocation.

Table 7: Full per-arm descriptive statistics on SEALED_TEST (machine-derived from run items; source tableA_descriptive.csv).

dataset	arm	n	ok	budg	EM	F1	evRec	retr	LLM
hotpotqa	static	2863	2718	117.0	0.565	0.714	0.000	1.817	4.00
hotpotqa	flat	2863	2683	144.0	0.570	0.715	0.000	1.788	4.00
hotpotqa	chain	2863	2504	323.0	0.561	0.711	0.000	1.516	3.00
2wikimulti-hop	static	4931	4483	265.0	0.540	0.645	0.000	1.763	3.00
2wikimulti-hop	flat	4931	4502	246.0	0.504	0.631	0.000	1.774	3.00
2wikimulti-hop	chain	4931	4746	2.000	0.497	0.629	0.000	1.870	4.00
musique	static	867.0	774.0	37.000	0.291	0.420	0.000	1.439	4.00
musique	flat	867.0	827.0	0.000	0.285	0.415	0.000	1.612	4.00
musique	chain	867.0	830.0	0.000	0.287	0.419	0.000	1.537	4.00

Table 8: Paired chain-vs-static differences on shared questions, bootstrap 95% CI (seed 2027). Source tableB_paired.csv.

dataset	metric	chain	static	Δ	CI _{2.5}	CI _{97.5}
hotpotqa	EM	0.542	0.541	0.001	0.000	1.000
hotpotqa	retrieval_calls	1.803	2.003	-0.200	-3.000	0.000
hotpotqa	llm_calls	49.320	25.321	23.998	-147.9	176.4
hotpotqa	documents	8.993	9.993	-1.000	-15.000	0.000
hotpotqa	latency_s	268728	144761	123966	-556810	1197249
2wikimulti-hop	EM	0.478	0.517	-0.039	-1.000	1.000
2wikimulti-hop	retrieval_calls	1.802	1.792	0.010	-1.000	1.000
2wikimulti-hop	llm_calls	14.210	37.130	-22.920	-113.8	7.000
2wikimulti-hop	documents	8.876	8.826	0.049	-5.000	5.000
2wikimulti-hop	latency_s	63999	174855	-110856	-447180	30577
musique	EM	0.275	0.264	0.010	0.000	0.000
musique	retrieval_calls	1.472	1.563	-0.091	-3.000	1.000
musique	llm_calls	5.543	19.536	-13.993	-91.700	2.000
musique	documents	6.566	7.007	-0.441	-12.700	3.350
musique	latency_s	31648	143772	-112124	-633966	14347

Table 9: Execution status by arm $\times$ dataset (source tableC_failure_by_arm.csv).

dataset	arm	ok	budget	boundary	other	total	ok%
hotpotqa	static	2718	117.0	28.000	28.000	2863	94.900
hotpotqa	flat	2683	144.0	36.000	36.000	2863	93.700
hotpotqa	chain	2504	323.0	36.000	36.000	2863	87.500
2wikimulti-hop	static	4483	265.0	183.0	183.0	4931	90.900
2wikimulti-hop	flat	4502	246.0	183.0	183.0	4931	91.300
2wikimulti-hop	chain	4746	2.000	183.0	183.0	4931	96.200
musique	static	774.0	37.000	56.000	56.000	867.0	89.300
musique	flat	827.0	0.000	40.000	40.000	867.0	95.400
musique	chain	830.0	0.000	37.000	37.000	867.0	95.700

Table 10: Evidence-quality metrics by arm (source tableD_evidence.csv).

dataset	arm	evRecall	evRecall(med)	evMRR	docs
hotpotqa	static	0.000	0.000	0.000	9.060
hotpotqa	flat	0.000	0.000	0.000	8.914
hotpotqa	chain	0.000	0.000	0.000	7.555
2wikimulti-hop	static	0.000	0.000	0.000	8.690
2wikimulti-hop	flat	0.000	0.000	0.000	8.743
2wikimulti-hop	chain	0.000	0.000	0.000	9.211
musique	static	0.000	0.000	0.000	6.429
musique	flat	0.000	0.000	0.000	7.207
musique	chain	0.000	0.000	0.000	6.859

We pin this as a tested invariant of the system rather than an empirical correlation: the derivation is exact on the sub-domain for which the claim is made, and we do not assert

a learned estimator adds value on this subdomain. On non-well-defined chains the closed form is a proxy artifact and is outside the claim.

9.2 Logical-to-Physical Crossover (RQ9)

No physical implementation dominates universally, and on the full SEALED_TEST set the crossover is *heterogeneous* rather than uniformly favorable to the allocator. The chain law gives the requirement-aware arm genuine allocation freedom only where plans are deep; on the shallow plans that dominate 2WikiMultiHopQA the same freedom yields *no* retrieval saving and a significantly worse EM (Table 3, McNemar $p < 0.0001$). Mean realized retrieval calls are near-identical between chain and static on every dataset (paired Δ : HotpotQA -0.20 , 2Wiki $+0.01$, MuSiQue -0.09 calls/question), so the allocator does *not* spend fewer retrieval calls in aggregate; what it spends instead is LLM calls— $+24.0$ per question more than static on HotpotQA ($p < 0.0001$), -22.9 on 2Wiki, and -14.0 on MuSiQue. The crossover is therefore method-dependent: the chain arm trades LLM calls for a (mostly neutral, sometimes negative) EM effect, and it is strictly dominated on cost on 2WikiMultiHopQA where its EM is also significantly worse. We report this frontier without selecting the subdomain that shows a saving.

9.3 Bottleneck Attribution (RQ8)

Execution traces reveal where additional retrieval ceases to translate into answer gains. We decompose failures into acquisition, verification/materialization, packing, integration, and generation stages and attribute each incorrect or incomplete answer to a stage. The dominant bottleneck across the workloads is the *selection ceiling*—the point beyond which retrieving more evidence does not improve the answer—which is exactly why the matched-budget EM differences are small: on shallow plans the allocation is already near its floor, so the allocator has little to change, and on deep plans the extra exploration the chain arm buys is often spent on requirements that do not carry the answer. The honest complement to the effectiveness result is that the chain arm’s deeper exploration on the deep HotpotQA plans converts into *more* budget-exceeded questions (323 vs. 117 for static) rather than into more correct answers—a cost, not a benefit, on that dataset. The budget-exceeded total is not uniform across arms or datasets: on 2WikiMultiHopQA the static and flat arms hit the ceiling far more often (static 265, flat 246) than the chain arm (2), because the shallow 2Wiki plans exhaust their budget before the allocator’s deeper exploration can engage.

Figure 14 and Table 11 give the full status breakdown across all 25,983 executions. Of the 7,916 non-ok executions, 1,134 (4.4%) hit the matched retrieval budget—but the budget hits are *not* uniformly chain-skewed: they concentrate on the deep HotpotQA chain arm (323) and on the shallow 2Wiki static/flat arms (265/246), whereas the 2Wiki chain arm hits the ceiling only twice (2). The remaining 760 (2.9%)




Figure 14: SEALED_TEST status breakdown by dataset (chain arm). Completed, budget-exceeded, deterministic plan-question mismatch, and residual infrastructure failure. Non-completions are dominated by structural limits (budget, compiler-rejected mismatch), not by flaky serving.




Figure 15: RQ8: budget-exceeded count by arm $\times$ dataset. The budget hits are *not* uniformly chain-skewed: the deep HotpotQA chain arm dominates its own budget hits (323), but on 2WikiMultiHopQA the static (265) and flat (246) arms hit the ceiling far more often than the chain arm (2).

Table 11: Execution status across all 25,983 SEALED_TEST executions (8,661 questions $\times$ 3 arms).

status	count	% of total	note
completed (ok)	24,067	92.6	executed within budget
budget-exceeded	1,134	4.4	HotpotQA chain 323 + 2Wiki
deterministic boundary	760	2.9	compiler rejects: ANSWERUNREACHABLE
infrastructure residual	22	0.08	gateway-side provider errors
total	25,983	100.0	

are deterministic plan-question mismatches the compiler correctly rejects (ANSWERUNREACHABLE 732, no-join-path 16, dependency-cycle 9, validation-error 3). Only 22 (0.08%) are residual infrastructure failures (gateway-side provider errors). On the solvable subset (25,201 questions, excluding infrastructure and deterministic-boundary failures) the completed rate is 95.5%. The failure mix is itself a finding: the system’s non-completions are dominated by *structural* limits (budget, plan-question mismatch), not by flaky infrastructure, so the EM differences we report are not an artifact of unstable serving.

9.4 Runtime Trace Cases (RQ8)

The per-dataset frontier is reported in Table 3 and Section 8. On HotpotQA the chain arm issues +24.0 LLM calls/question more than static (paired CI[+21.7, +26.3], $p < 0.0001$) while realized retrieval calls are only marginally lower (−0.20); on 2Wiki the chain arm is both more expensive in LLM calls and significantly worse in EM; on MuSiQue it is cheaper in LLM calls (−14.0) and marginally best in EM (27.5%, McNemar $p=0.064$). The trace cases confirm the mechanism: the chain arm’s allocation freedom is real, but on shallow plans it is spent on requirements that do not move the answer, while on deep MuSiQue plans it occasionally helps. We do not report a uniform cost saving because none exists at the full-sample scale.

9.5 Scaling and Overhead (RQ9)

We quantify optimizer overhead as plan complexity increases, separating compile time, estimator inference, plan enumeration/pruning, and evidence execution. The optimizer compilation contributes ≈ 0 ms per query (the search is deterministic and bounded by ≤ 256 candidate orders $\times 2$ implementations per requirement); the overhead is concentrated in execution time, which scales with the number of retrieval and LLM calls the plan triggers. Optimizer search telemetry is auditable per query (candidates enumerated, frontier pruned, validation warnings recorded in the execution profile). Because the cost difference between arms is dominated by LLM-call count (not retrieval-call count) on the full sample, we report LLM invocations alongside retrieval calls as the primary cost axes in Table 3’s companion statistics, and we note that the chain arm’s higher HotpotQA LLM cost is the mechanical price of its deeper evidence-chain exploration under a hard budget.

9.6 Robustness to Estimator Error

Because the importance estimator is construction-derived on the claimed subdomain, it is not subject to calibration error there; its behavior is exact. We therefore report robustness as a structural property on well-defined chains, and note that generalization to non-well-defined chains would require a calibrated estimator, which is future work.

10 Conclusion

We formulated complex multi-hop retrieval as *declarative evidence execution*: a question is compiled into typed, dependent evidence requirements; a deterministic importance-weighted allocator assigns legal orders, physical implementations, and budget allocations to those requirements under a matched budget; and a deterministic executor acquires evidence with hard budget enforcement and full provenance. We introduced a typed evidence algebra whose state-transition semantics (`apply_observation`) make requirement progress first-class and auditable, and we derived a chain law pinning requirement importance to dependency depth ($\tau = 2d - 1$) on well-defined chains, so the estimator is construction-derived rather than learned.

The evaluation does *not* establish a universal accuracy or cost win, and we are deliberate about that boundary. On the frozen SEALED_TEST set (8,661 questions; model `qwen3.5-9b`; matched retrieval budget = 8) the chain-rule allocator is significantly *worse* than static on 2WikiMultiHopQA (EM 47.8% vs. 51.7%, paired bootstrap $\Delta\text{EM} -3.9\text{pt}$ CI[−5.0, −2.9], McNemar $p < 0.0001$), statistically tied on HotpotQA (flat 55.4% / static 54.1% / chain 54.2%; chain-vs-static McNemar $p=0.80$), and only marginally best on MuSiQue (chain 27.5%, McNemar $p=0.064$). The flat arm—which performs no dependency-tree search—is strongest or tied on every dataset. The matched-budget frontier is similarly heterogeneous: mean retrieval calls are near-identical between chain and static (HotpotQA −0.20, 2Wiki +0.01, MuSiQue −0.09 calls/question), while the chain arm spends more LLM calls on HotpotQA (+24.0 per question, $p < 0.0001$). The honest statement is that allocator benefit is conditional on dependency depth: it helps on the deep chains of MuSiQue and is neutral-to-harmful on the shallow plans that dominate 2WikiMultiHopQA, where the chain law leaves little allocation freedom. We report the absence of a global win with the same rigor as a win would demand. We treat runtime re-optimization over alternative physical plans as future work under a branching execution model, because on the chain-only topology this system runs today it is vacuous.

More broadly, this study suggests that complex retrieval can be treated as an execution problem in which evidence goals remain stable while physical acquisition strategies adapt to observed state and resource constraints. The honest path forward is to make the evidence requirements and their satisfaction state the interface between declarative programs and physical optimization, and to evaluate every claim—including the absence of a benefit—against a matched budget with the same rigor as the presence of one.

Limitations. The matched-budget headline numbers rest on the frozen SEALED_TEST set ($n=2,863/4,931/867$ HotpotQA / 2WikiMultiHop / MuSiQue, pooled $n=8,661$) under one decoder (`qwen3.5-9b` at temperature 0) through a shared provider; generalizability to other decoders is untested. Of the 25,983 executions, 24,067 completed (92.6%); on the solvable subset (25,201, excluding 22 residual infrastructure failures and 760 deterministic plan-question mismatches the compiler correctly rejects) the completed rate is 95.5%, and 1,134 questions exceeded the matched retrieval budget. The budget hits are not uniformly chain-skewed: they concentrate on the deep HotpotQA chain arm (323) and on the shallow 2Wiki static/flat arms (265/246), whereas the 2Wiki chain arm hits the ceiling only twice (2). Evidence corpora are the public HotpotQA/2Wiki/MuSiQue passage indices; the heterogeneous-evidence result (HoVer/FEVEROUS pilot) is a small transfer proof-of-concept, not a Wikipedia-scale retrieval benchmark. The chain law $\tau = 2d - 1$ is exact on well-defined *chains* where the compiled requirement order respects topological dependencies; on general (branching) graphs the enumerate-index proxy is a heuristic and is outside the claim’s scope. Cost accounting covers realized retrieval

and LLM calls under the matched budget; latency and monetary cost are reported for the optimizer overhead but are not the primary comparison axis.

References

- [1] SlotRAG-X Authors. 2026. Slot-based Retrieval-Augmented Generation Under a Matched Budget: An Honest Evaluation of Coverage Ceilings. *arXiv preprint*.
- [2] Y. Chen et al. 2026. Retrieval is Cheap, Show Me the Code: Executable Multi-Hop Reasoning for RAG. In *arXiv preprint arXiv:2605.12975*. arXiv:2605.12975
- [3] X. Liu et al. 2026. When RAG Meets Query Planning: Logical Query Trees for Resolving Exploratory Reasoning Problems. In *arXiv preprint arXiv:2607.00508*. arXiv:2607.00508
- [4] J. Wang et al. 2026. DynaKRAG: A Unified Framework for Learnable Evidence Control in Multi-Hop RAG. In *arXiv preprint arXiv:2607.06507*. arXiv:2607.06507